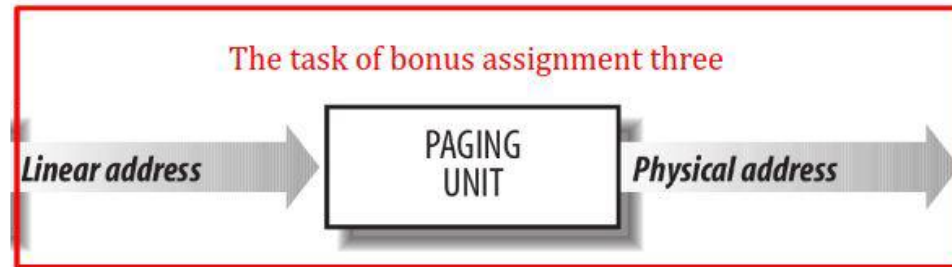


Bonus Assignment Four

XUE JIN

jinxue@cse.cuhk.edu.hk

Overview



Task:

- Implement a simple simulator of translating processes' linear address (virtual address) to physical address using paging technique.

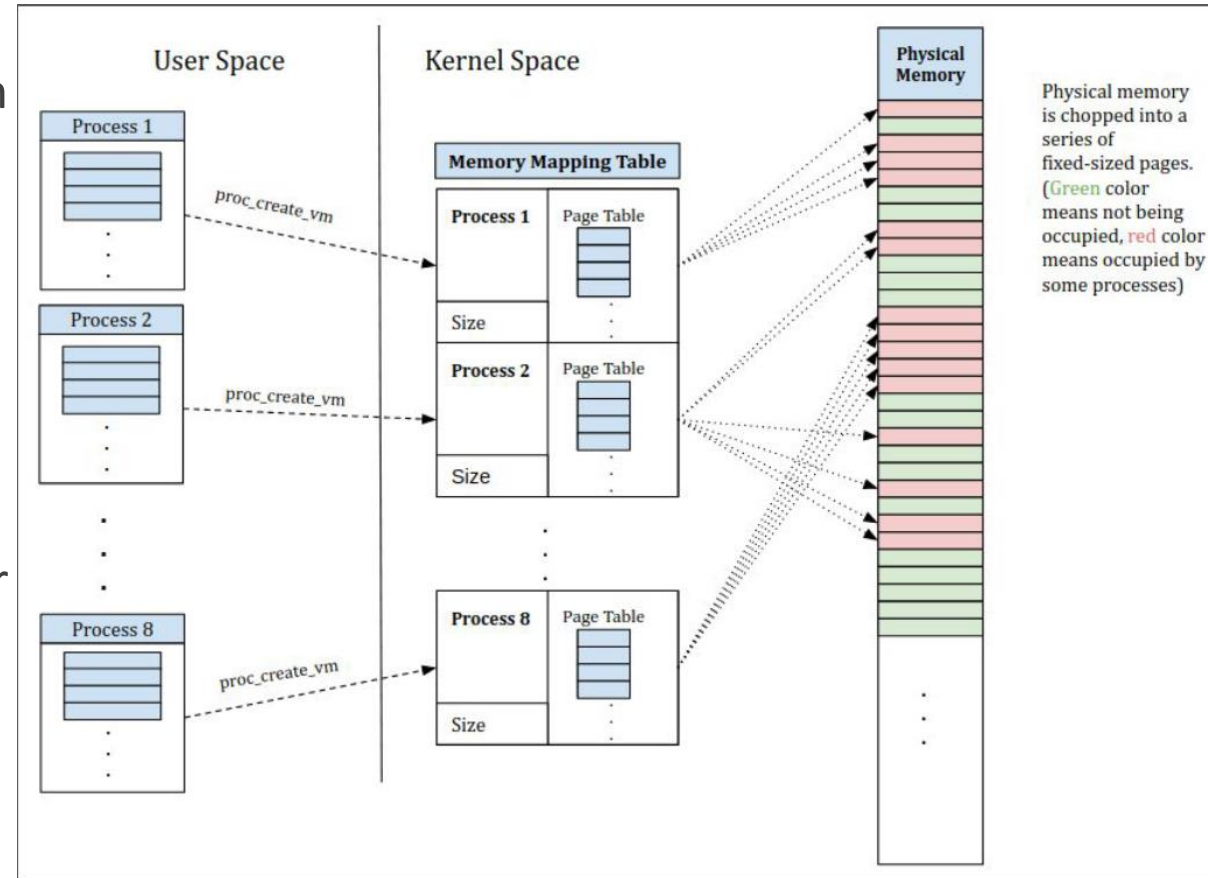
Overview

User Space

- User can **create** (`proc_create_vm`) process with specified virtual memory size through the API.
- User can **read/write** (`vm_read/vm_write`) the process' virtual memory through API.
- User can **terminate** (`proc_exit_vm`) process through the API.

Kernel Space

- The kernel simulator maintains a **page table** for each process and a **shared physical memory**.
- The kernel simulator handles the mappings from processes' virtual memory to physical memory.



Data Structures

Kernel

- **space**: the physical memory.
- **allocated_pages**: Total number of allocated pages for all running processes.
- **occupied_pages**: char array (size is the number of kernel pages) to indicate the free pages, 0 -> free, 1 -> occupied.
- **running**: an array marking if the corresponding process is running.
- **mm**: an array of page tables.
 - Array of **MMStruct**

```
/*
Reference: textbook chapter 18.
To make it simple, we do not encode PFN and flag bits together to an integer.
PTE: page table entry.
PFN: page frame number (here we use it to indicate the page id in kernel managed memory).

present: represents if this translation is built, 0 -> not built, 1 -> built.
Currently when the pages are allocated (proc_create_vm), PFN will be set to -1 and
present will be set to 0 because the translation is not yet built.
After you access this page (vm_read && vm_write), you will need to build the translation and present will be set to 1.
*/
struct PTE {
    int PFN;
    char present;
};

struct PageTable {
    struct PTE* ptes;
};

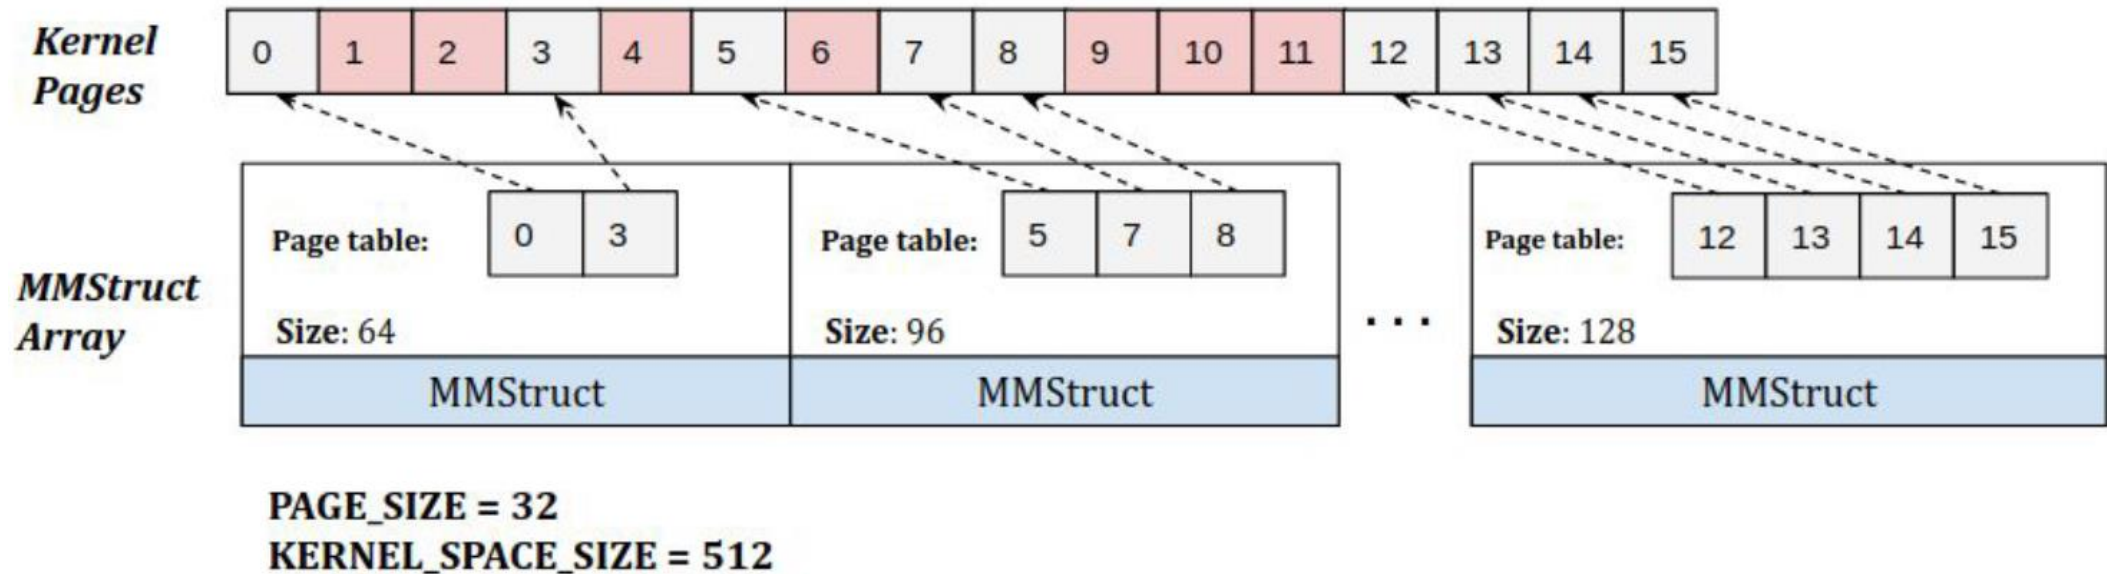
/*
1. The user space and the user space page id start from 0.
2. size indicates the size of user space (&& kernel-managed memory) allocated for this process.
3. page_table is an array of PTE (page table entry).
*/
struct MMStruct {
    int size;
    struct PageTable* page_table;
};

// The Kernel manages MAX_PROCESS_NUM of processes.
struct Kernel {
    char* space;
    int allocated_pages; // The number of allocated pages for processes.
    char* occupied_pages; // For simplicity, we use a char array to indicate the free pages, 0 for free, 1 for occupied.
    char* running; // An array marking if the process is running.
    struct MMStruct* mm; // An array of MMStruct for each process.
};
```

Data Structures (Cont'd)

MMStruct:

- **size** determines the number of allocated pages.
- **page table** contains an array of PTEs.



Implementation

Task: Totally 4 functions (API) in the kernel.c need to be implemented.

1. `int proc_create_vm(struct Kernel* kernel, int size)`

- 1.1. Check if a free process slot exists (check the running, the slot will be the pid returned).
- 1.2. Check if there's enough free space (check `allocated_pages`).
- 1.3. Allocate the space for `page_table` (the size of it depends on the pages you needed. e.g. if `size=33` and `PAGE_SIZE=32`, then you need 2 pages) and update `allocated_pages`.
- 1.4. The mapping to kernel-managed memory is not built up, all the PFN should be set to -1 and present byte to 0 (PTE) and set the corresponding element in `running` to 1.
- 1.5. Return the pid if success, -1 if failure.

Implementation (Cont'd)

2. `int vm_read(struct Kernel* kernel, int pid, char* addr, int size, char* buf)`

2.1. Check if the reading range is out-of-bounds.

2.2. If the pages in the range [addr, addr+size) of the user space of that process are not present, you should firstly map them to the free kernel-managed memory pages (**first fit policy: scan from the beginning**).

3. `int vm_write(struct Kernel* kernel, int pid, char* addr, int size, char* buf)`

4. `int proc_exit_vm(struct Kernel* kernel, int pid)`

4.1. Reset the corresponding pages in `occupied_pages` to 0.

4.2. Release the `page_table` in the corresponding `MMStruct` and set to NULL. Return 0 when success, -1 when failure.

Demo

You can find the demo program in *demo.c*

Initial parameters

```
----- Demo Program -----  
KERNEL_SPACE_SIZE=8192  
VIRTUAL_SPACE_SIZE=512  
PAGE_SIZE=32  
MAX_PROCESS_NUM=8
```

Create process 0 and 1

```
// Create process 0 with size VIRTUAL_SPACE_SIZE.  
int pid0 = proc_create_vm(kernel, VIRTUAL_SPACE_SIZE);  
if(pid0 == 0)  
    score += 1;  
  
// Create process 1 with size VIRTUAL_SPACE_SIZE/2.  
int pid1 = proc_create_vm(kernel, VIRTUAL_SPACE_SIZE/2);  
if(pid1 == 1)  
    score += 1;
```

```
Before reading pages 0-7 of process 1  
free space: (addr:0, size:8192)  
Memory mappings of process 1  
virtual page 0: Not present  
virtual page 1: Not present  
virtual page 2: Not present  
virtual page 3: Not present  
virtual page 4: Not present  
virtual page 5: Not present  
virtual page 6: Not present  
virtual page 7: Not present
```


Demo (Cont'd)

Read page 0-7 of process 1

```
// Check the free space after reading pages 0-7 for process 1.
memset(buf, 0, 128);
memset(temp_buf, 0, 512);
vm_read(kernel, pid1, (char *)0, 234, temp_buf);
get_kernel_free_space_info(kernel, buf);
if(strcmp(buf, "free space: (addr:256, size:7936)\n") == 0)
    score += 1;
```

```
After reading pages 0-7 of process 1
free space: (addr:256, size:7936)
Memory mappings of process 1
virtual page 0 -> physical page 0
virtual page 1 -> physical page 1
virtual page 2 -> physical page 2
virtual page 3 -> physical page 3
virtual page 4 -> physical page 4
virtual page 5 -> physical page 5
virtual page 6 -> physical page 6
virtual page 7 -> physical page 7
```

Create process 2

```
// Create process 2 with size VIRTUAL_SPACE_SIZE/4.
int pid2 = proc_create_vm(kernel, VIRTUAL_SPACE_SIZE/4);
if(pid2 == 2)
    score += 1;
```

```
Before writing/reading page 1 of process 2
free space: (addr:256, size:7936)
Memory mappings of process 2
virtual page 0: Not present
virtual page 1: Not present
virtual page 2: Not present
virtual page 3: Not present
```

Demo (Cont'd)

Write and read the page 1 of process 2

```
// Verify writting/reading page 1 for process 2.
memset(buf, 0, 128);
memset(temp_buf, 0, 512);
temp_buf[0] = 'a';
if(vm_write(kernel, pid2, (char*)(42), 1, temp_buf) == 0)
    score += 1;
temp_buf[0] = '\0'; // Clean the first byte of temp_buf.
if(vm_read(kernel, pid2, (char*)(42), 1, temp_buf) == 0)
    score += 1;
if(temp_buf[0] == 'a')
    score += 1;
get_kernel_free_space_info(kernel, buf);
if(strcmp(buf, "free space: (addr:288, size:7904)\n") == 0)
    score += 1;
```

```
After writting/reading page 1 of process 2
free space: (addr:288, size:7904)
Memory mappings of process 2
virtual page 0: Not present
virtual page 1 -> physical page 8
virtual page 2: Not present
virtual page 3: Not present
```

Demo (Cont'd)

Create process 3 and write page 0-3 of process 3

```
After writting pages 0-3 of process 3
free space: (addr:416, size:7776)
Memory mappings of process 3
virtual page 0 -> physical page 9
virtual page 1 -> physical page 10
virtual page 2 -> physical page 11
virtual page 3 -> physical page 12
```

Exit

```
After process 2 exits
free space: (addr:256, size:32) -> (addr:416, size:7776)
After process 3 exits
free space: (addr:256, size:7936)
After process 1 exits
free space: (addr:0, size:8192)
After process 0 exits
free space: (addr:0, size:8192)
Full Score: 19, Your Score: 19
```

Submission (Deadline: 23:59, May 8, 2022)

After you finish bonus assignment four, please **only submit the kernel.c** to Blackboard.